

DSC 630 Term Project

Christina Vosnak

February 9th, 2025

Impact of Air Quality Metrics and Weather Conditions on Health

```
In [17]: # Import Libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [2]: # Load data
data = pd.read_csv('air_quality_health_impact_data.csv')
data.head()
```

```
Out[2]:
```

	RecordID	AQI	PM10	PM2_5	NO2	SO2	O3	Temp
0	1	187.270059	295.853039	13.038560	6.639263	66.161150	54.624280	5.
1	2	475.357153	246.254703	9.984497	16.318326	90.499523	169.621728	1.
2	3	365.996971	84.443191	23.111340	96.317811	17.875850	9.006794	1.
3	4	299.329242	21.020609	14.273403	81.234403	48.323616	93.161033	21.
4	5	78.009320	16.987667	152.111623	121.235461	90.866167	241.795138	9.

```
In [3]: # Check Size of the data
data.shape
```

```
Out[3]: (5811, 15)
```

```
In [4]: # Check data types for each column
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5811 entries, 0 to 5810
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RecordID              5811 non-null   int64
1   AQI                   5811 non-null   float64
2   PM10                  5811 non-null   float64
3   PM2_5                 5811 non-null   float64
4   NO2                   5811 non-null   float64
5   SO2                   5811 non-null   float64
6   O3                    5811 non-null   float64
7   Temperature           5811 non-null   float64
8   Humidity              5811 non-null   float64
9   WindSpeed             5811 non-null   float64
10  RespiratoryCases      5811 non-null   int64
11  CardiovascularCases   5811 non-null   int64
12  HospitalAdmissions    5811 non-null   int64
13  HealthImpactScore     5811 non-null   float64
14  HealthImpactClass     5811 non-null   float64
dtypes: float64(11), int64(4)
memory usage: 681.1 KB

```

```

In [5]: # Check for missing values
data.isnull().sum()

```

```

Out[5]: RecordID          0
        AQI              0
        PM10            0
        PM2_5           0
        NO2             0
        SO2             0
        O3              0
        Temperature     0
        Humidity         0
        WindSpeed        0
        RespiratoryCases 0
        CardiovascularCases 0
        HospitalAdmissions 0
        HealthImpactScore 0
        HealthImpactClass 0
dtype: int64

```

```

In [6]: # Check for duplicate values
data.duplicated().sum()

```

```

Out[6]: 0

```

The information above shows that there is no missing data in the dataset.

```

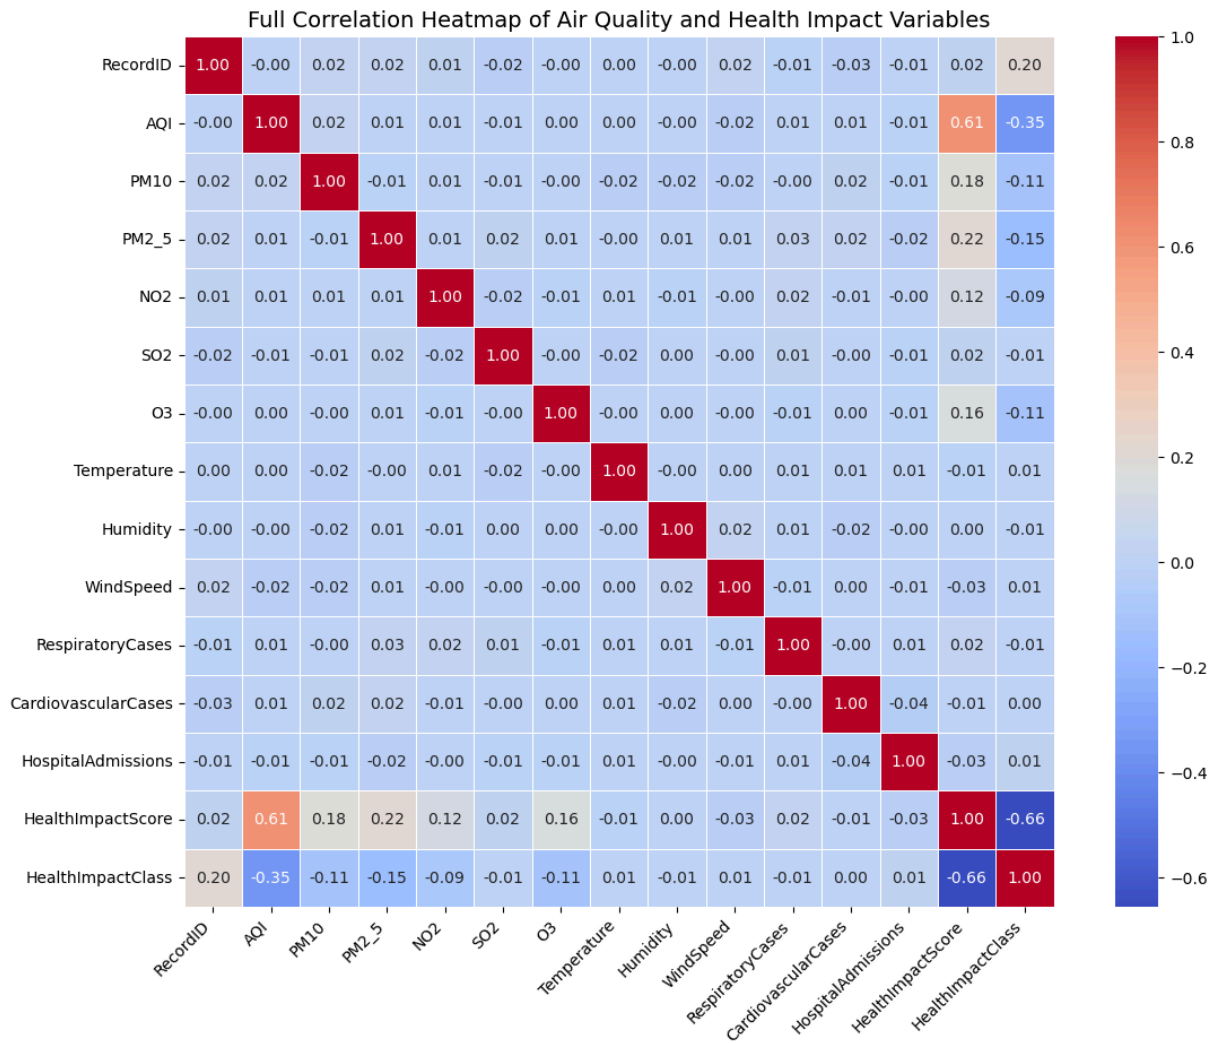
In [7]: # Create Correlation Heatmap to conduct preliminary analysis
plt.figure(figsize=(14, 10))
heatmap = sns.heatmap(data.corr(), annot=True, cmap='coolwarm', fmt=".2f", linewidths=1)

# Improve layout and readability

```

```
plt.title("Full Correlation Heatmap of Air Quality and Health Impact Variables", fo
plt.xticks(rotation=45, ha="right", fontsize=10)
plt.yticks(fontsize=10)

plt.show()
```




```
In [8]: # Create Chart Comparing AQI and Health Impact Score
plt.figure(figsize=(10, 6))

# Scatter plot comparing AQI and Health Impact Score
sns.scatterplot(x=data['AQI'], y=data['HealthImpactScore'], alpha=0.5)

# Titles and Labels
plt.title("Comparison of AQI and Health Impact Score", fontsize=14)
plt.xlabel("Air Quality Index (AQI)", fontsize=12)
plt.ylabel("Health Impact Score", fontsize=12)

# Show the plot
plt.show()
```




```
In [9]: # Create Chart Comparing PM10 and Health Impact Score
plt.figure(figsize=(10, 6))

# Scatter plot comparing AQI and Health Impact Score
sns.scatterplot(x=data['PM10'], y=data['HealthImpactScore'], alpha=0.5)

# Titles and Label
plt.title("Comparison of PM10 and Health Impact Score", fontsize=14)
plt.xlabel("PM10", fontsize=12)
plt.ylabel("Health Impact Score", fontsize=12)

# Show the plot
plt.show()
```




```
In [10]: # Create Chart Comparing PM2_5 and Health Impact Score
plt.figure(figsize=(10, 6))

# Scatter plot comparing PM2_5 and Health Impact Score
sns.scatterplot(x=data['PM2_5'], y=data['HealthImpactScore'], alpha=0.5)

# Titles and Labels
plt.title("Comparison of PM2_5 and Health Impact Score", fontsize=14)
plt.xlabel("PM2_5", fontsize=12)
plt.ylabel("Health Impact Score", fontsize=12)

# Show the plot
plt.show()
```



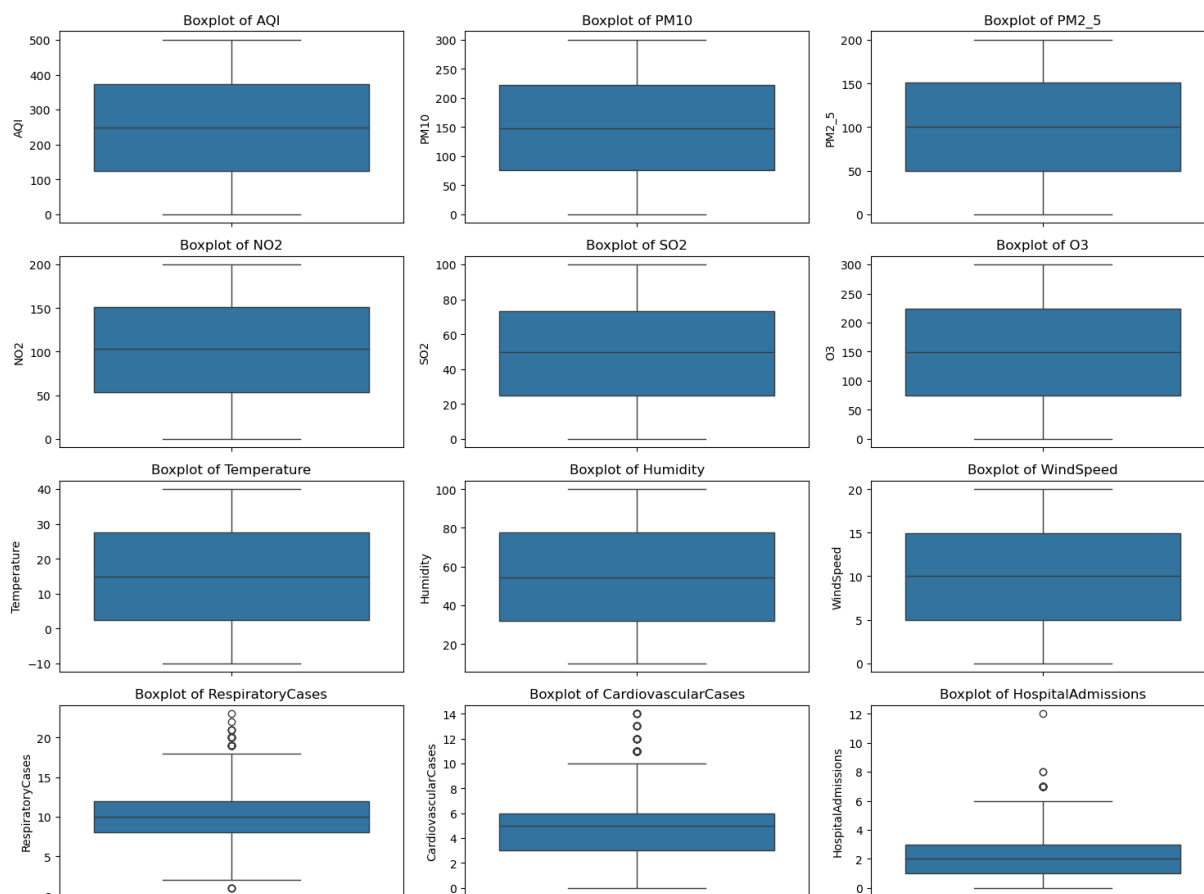
```
In [11]: # Creating boxplots for outlier detection in key variables
fig, axes = plt.subplots(4, 3, figsize=(15, 12))
fig.suptitle("Boxplots for Outlier Detection in Air Quality and Health Impact Varia

# List of columns to plot
columns_to_plot = ['AQI', 'PM10', 'PM2_5', 'NO2', 'SO2', 'O3', 'Temperature', 'Humi

# Create boxplots for each variable
for ax, col in zip(axes.flatten(), columns_to_plot):
    sns.boxplot(y=data[col], ax=ax)
    ax.set_title(f"Boxplot of {col}")

# Adjust layout for better readability
plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()
```

Boxplots for Outlier Detection in Air Quality and Health Impact Variables



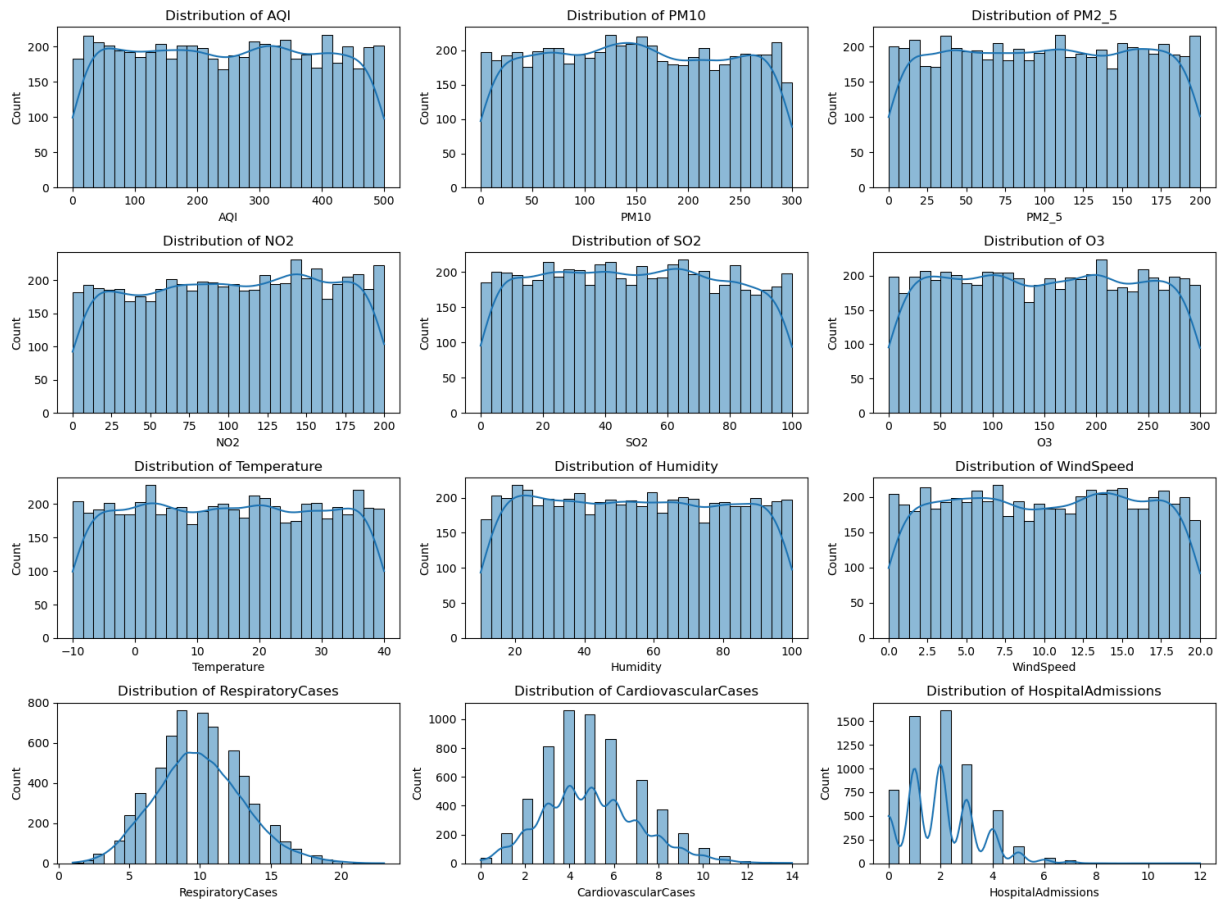
```
In [12]: # Creating distribution plots for key air quality and health impact variables
fig, axes = plt.subplots(4, 3, figsize=(15, 12))
fig.suptitle("Distribution of Air Quality and Health Impact Variables", fontsize=16)

# List of columns to plot
columns_to_plot = ['AQI', 'PM10', 'PM2_5', 'NO2', 'SO2', 'O3', 'Temperature', 'Humidity', 'WindSpeed', 'RespiratoryCases', 'CardiovascularCases', 'HospitalAdmissions']

# Create distribution plots for each variable
for ax, col in zip(axes.flatten(), columns_to_plot):
    sns.histplot(data[col], bins=30, kde=True, ax=ax)
    ax.set_title(f"Distribution of {col}")

# Adjust layout for better readability
plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()
```

Distribution of Air Quality and Health Impact Variables



```
In [13]: # Function to count outliers using IQR
def count_outliers(column, data):
    Q1 = data[column].quantile(0.25)
    Q3 = data[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = data[column][(data[column] < lower_bound) | (data[column] > upper_bound)]
    return len(outliers)

# Detect outliers for all numerical features
numerical_features = data.select_dtypes(include=['float64', 'int64']).columns
outliers = []
for col in numerical_features:
    outliers.append({
        'Feature': col,
        'Num of outliers': count_outliers(col, data)
    })

# Convert to DataFrame and display
outliers_df = pd.DataFrame(outliers)
print(outliers_df)
```


	Feature	Num of outliers
0	RecordID	0
1	AQI	0
2	PM10	0
3	PM2_5	0
4	NO2	0
5	SO2	0
6	O3	0
7	Temperature	0
8	Humidity	0
9	WindSpeed	0
10	RespiratoryCases	42
11	CardiovascularCases	74
12	HospitalAdmissions	31
13	HealthImpactScore	1352
14	HealthImpactClass	1003

```
In [14]: # Create Features and Target Variable
X = data.drop(columns=['HealthImpactScore', 'HealthImpactClass'])
y = data['HealthImpactScore']
```

```
In [18]: # One-hot encode categorical variables
X = pd.get_dummies(X, drop_first=True)

# Standardize numerical features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [19]: X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, ran
```

```
In [20]: lr = LinearRegression()
lr.fit(X_train, y_train)
```

```
Out[20]: ▾ LinearRegression
LinearRegression()
```

```
In [21]: rf = RandomForestRegressor(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)
```

```
Out[21]: ▾ RandomForestRegressor
RandomForestRegressor(random_state=42)
```

```
In [23]: models = {"Linear Regression": lr, "Random Forest": rf}

for name, model in models.items():
    y_pred = model.predict(X_test)
    print(f"Model: {name}")
    print("MAE:", mean_absolute_error(y_test, y_pred))
    print("MSE:", mean_squared_error(y_test, y_pred))
    print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
```

```
print("R2 Score:", r2_score(y_test, y_pred))
print("-" * 40)
```

Model: Linear Regression
MAE: 7.356775305712758
MSE: 92.6651461558134
RMSE: 9.626273741994531
R² Score: 0.5054324194576925

Model: Random Forest
MAE: 1.5758572616603441
MSE: 10.329369986851852
RMSE: 3.213933724713665
R² Score: 0.9448706257438612

```
In [24]: # Extract feature importances from the trained Random Forest model
feature_importances = rf.feature_importances_

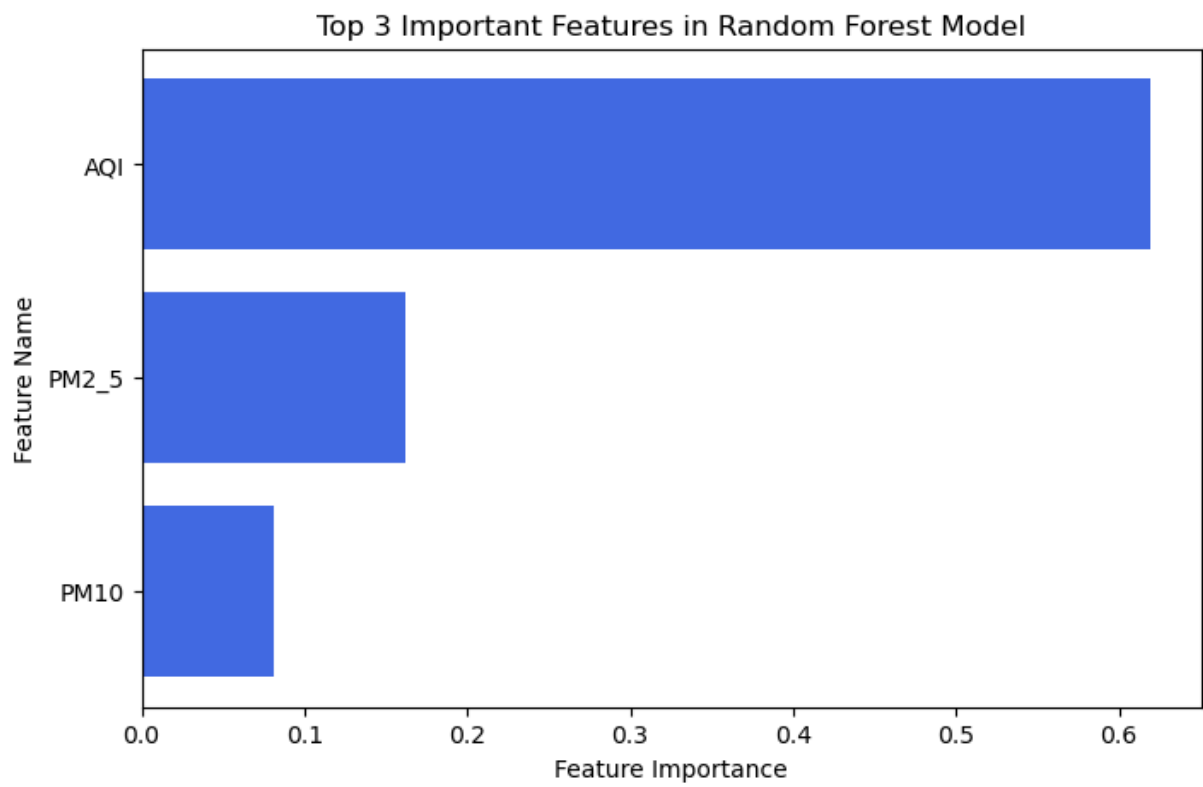
# Create a DataFrame to pair feature names with their importance scores
feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importances
})

# Sort features by importance in descending order
top_features = feature_importance_df.sort_values(by="Importance", ascending=False).

# Display the top 3 features
print(top_features)

# Plot the top 3 features
plt.figure(figsize=(8, 5))
plt.barh(top_features["Feature"], top_features["Importance"], color="royalblue")
plt.xlabel("Feature Importance")
plt.ylabel("Feature Name")
plt.title("Top 3 Important Features in Random Forest Model")
plt.gca().invert_yaxis() # Invert y-axis for better visualization
plt.show()
```

	Feature	Importance
1	AQI	0.619246
3	PM2_5	0.161854
2	PM10	0.080753



In []: